

---

# Latent Reasoning VLA: Latent Thinking and Prediction for Vision-Language-Action Models

---

000  
001  
002  
003  
004  
005  
006  
007  
008  
009  
010  
011  
012  
013  
014  
015  
016  
017  
018  
019  
020  
021  
022  
023  
024  
025  
026  
027  
028  
029  
030  
031  
032  
033  
034  
035  
036  
037  
038  
039  
040  
041  
042  
043  
044  
045  
046  
047  
048  
049  
050  
051  
052  
053  
054  
Anonymous Authors<sup>1</sup>

## Abstract

Vision-Language-Action (VLA) models benefit from chain-of-thought (CoT) reasoning, but existing approaches incur high inference overhead and rely on discrete reasoning representations that mismatch continuous perception and control. We propose Latent Reasoning VLA (**LaRA-VLA**), a unified VLA framework that internalizes multimodal CoT reasoning into continuous latent representations for embodied action. LaRA-VLA performs unified reasoning and prediction in latent space, eliminating explicit CoT generation at inference time and enabling efficient, action-oriented control. To realize latent embodied reasoning, we introduce a curriculum-based training paradigm that progressively transitions from explicit textual and visual CoT supervision to latent reasoning, and finally adapts latent reasoning dynamics to condition action generation. We construct two structured CoT datasets and evaluate LaRA-VLA on both simulation benchmarks and long-horizon real-robot manipulation tasks. Experimental results show that LaRA-VLA consistently outperforms state-of-the-art VLA methods while reducing inference latency by up to 90% compared to explicit CoT-based approaches, demonstrating latent reasoning as an effective and efficient paradigm for real-time embodied control.

## 1. Introduction

Vision-Language-Action (VLA) models have emerged as a promising direction for scalable, general-purpose robotic manipulation (Kim et al., 2025b), as they aim to end-to-end map rich multimodal observations and language instructions to continuous control commands. To improve VLA performance, a variety of training and inference strategies




Figure 1. Comparison of CoT formulations in VLA models. (a) Textual CoT-based VLA explicitly generates discrete textual reasoning tokens, which are subsequently decoded into actions via downstream modules such as autoregressive (AR) policies or action experts (AE). (b) Visual CoT-Based VLA represents reasoning through discrete visual goal tokens, obtained via detokenization or alignment, before action generation. (c) Our model internalizes both textual and visual reasoning into continuous latent representations. Visual goal latents are aligned with perceptual features and serve a dual role: they encode future-oriented task information and provide implicit supervision for textual CoT latents.

have been proposed (Bai et al., 2025b; Sridhar et al., 2025; Lin et al., 2025b). Among these approaches, incorporating Chain-of-Thought (CoT) reasoning has proven particularly effective (Chen et al., 2025c). CoT is commonly framed as a teaching signal that encourages structured intermediate reasoning, thereby improving interpretability and generalization in complex, multi-step robotic tasks.

Existing CoT approaches for VLA models can be broadly categorized by modality, as shown in Figure 1. Text-based CoT methods represent intermediate reasoning explicitly in natural language, covering task decomposition and high-level planning, and may additionally verbalize information

<sup>1</sup>Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

055 derived from visual observations (Zawalski et al., 2025; Sun  
 056 et al., 2025b; Huang et al., 2025). Visual CoT methods  
 057 instead express reasoning through visual reconstruction or  
 058 prediction, typically by modeling future observations or  
 059 intermediate visual states (Zhao et al., 2025; Zhang et al.,  
 060 2025b). A third line of work combines textual and visual  
 061 CoT, leveraging multi-modal intermediate representations  
 062 for reasoning (Zhang et al., 2025a).

063 Despite their effectiveness, existing CoT-based methods  
 064 face two fundamental challenges. First, text-based CoT often  
 065 requires extensive textual reasoning during inference,  
 066 leading to substantial computational overhead. Long rea-  
 067 soning traces dramatically increase token length, resulting  
 068 in excessive KV-cache usage, high memory consumption,  
 069 and increased latency. In practice, such models may operate  
 070 at control frequencies below 5 Hz, or even around 1 Hz (Za-  
 071 walski et al., 2025), which is unacceptable for real-time  
 072 robotic control. Second, most CoT formulations rely on  
 073 discrete representations: textual CoT is constrained to lan-  
 074 guage tokens (Zawalski et al., 2025; Lin et al., 2025a), while  
 075 visual CoT is typically aligned with discrete visual tokens  
 076 produced by VQ-based tokenizers (Zhao et al., 2025). We  
 077 argue that CoT is effective not because it is expressed in nat-  
 078 ural language, but because it exposes structured intermediate  
 079 reasoning. In embodied settings, where both perception and  
 080 action evolve in continuous spaces, constraining reasoning  
 081 to discrete tokens introduces a representational mismatch  
 082 between reasoning and control.

083 To address these challenges, we propose Latent Reason-  
 084 ing VLA (LaRA-VLA), a unified latent-reasoning VLA  
 085 framework that performs reasoning and prediction entirely  
 086 in latent space for robotic action. A comparison between  
 087 LaRA-VLA and existing CoT-based VLA methods is sum-  
 088 marized in Table 1. Inspired by latent chain-of-thought rea-  
 089 soning methods (Hao et al., 2024; Pham & Ngo, 2025), we  
 090 adapt latent reasoning to embodied VLA models, enabling  
 091 structured reasoning to be internalized into continuous latent  
 092 representations. To support this paradigm, we construct an  
 093 automated CoT annotation pipeline that provides structured  
 094 supervision, including subtask decomposition, movement  
 095 reasoning, and target object localization, and curate latent  
 096 reasoning datasets across both simulated and real-world en-  
 097 vironments, including LIBERO-LaRA, Bridge-LaRA, and  
 098 real-robot demonstrations.

099 Our training paradigm follows a three-stage progression that  
 100 transitions from explicit multi-modal reasoning to latent em-  
 101 bodied reasoning. Training begins with unified textual and  
 102 visual CoT supervision, where textual reasoning steps and  
 103 future visual predictions are explicitly aligned with their cor-  
 104 responding annotations. We then adopt a curriculum-based  
 105 strategy that models embodied reasoning as a sequence of  
 106 continuous latent states, gradually internalizing reasoning

107 structure by replacing explicit textual CoT with a compact  
 108 set of latent representations and relying on visual prediction  
 109 objectives as implicit supervision. Once textual CoT is fully  
 110 absorbed into latent reasoning, the model is further adapted  
 111 to action generation, enabling latent reasoning dynamics to  
 112 directly condition continuous control outputs. Throughout  
 113 training, visual latents used for supervision are encoded by  
 114 the same visual encoder as the input observations, and an  
 115 exponential moving average (EMA) encoder is employed as  
 116 a target network to stabilize latent representation learning  
 117 and prevent representation collapse.

118 Overall, our approach extends latent CoT reasoning from  
 119 language-only settings to embodied VLA models, enabling  
 120 efficient, action-oriented reasoning without explicit CoT  
 121 generation at inference time. By coupling latent reasoning  
 122 dynamics directly with action generation, LaRA-VLA allo-  
 123 cates computation to compact latent “thought steps” rather  
 124 than verbose textual reasoning, substantially reducing token  
 125 expansion and inference latency. This makes the approach  
 126 well-suited for real-time robotic control. Moreover, repre-  
 127 senting reasoning as continuous latent dynamics avoids the  
 128 representational mismatch introduced by discrete language  
 129 or visual tokens. Our contributions are threefold:

- We introduce a latent-reasoning paradigm for Vision–Language–Action models, in which chain-of-thought reasoning is internalized into continuous latent representations across textual and visual modalities, enabling inference-efficient reasoning that aligns with continuous perception and control.
- We propose LaRA-VLA, a unified VLA model that realizes this paradigm through a curriculum-based training strategy, progressively transitioning from explicit multi-modal CoT supervision to latent embodied reasoning, guided by predictive visual latent objectives and stabilized with EMA-based visual encoders.
- We construct two structured chain-of-thought datasets, LIBERO-LaRA and Bridge-LaRA, featuring multi-modal reasoning annotations for embodied manipulation, and conduct extensive evaluations in both simulation and long-horizon real-robot tasks to demonstrate the effectiveness and robustness of LaRA-VLA.

## 2. Related Work

### 2.1. Vision Language Action Models

Vision–Language–Action (VLA) models extend large multi-modal language models to robotic control, with extensive efforts devoted to architectural advances (Black et al., 2024; Intelligence et al., 2025; Kim et al., 2025a) as well as training- and inference-time optimizations (Sridhar et al.,

110  
 111 *Table 1.* A taxonomy of VLA models based on the representation forms of chain-of-thought CoT and actions. Specifically, we categorize  
 112 models by whether their textual CoT is represented as explicit discrete tokens or continuous latent states, whether their visual CoT aligns  
 113 with raw pixels, discrete visual tokens, or encoded continuous visual latents, and whether the action output is represented as discrete  
 114 tokens or continuous values.  
 115

Method	Venue	Text CoT		Visual CoT		Action
		Presence	Form	Presence	Align Form	Form
<i>VLAs with Text CoT</i>						
ECoT (Zawalski et al., 2025)	CoRL 2024	✓	Discrete Token	✗	—	Discrete
GraspVLA (Deng et al., 2025)	CoRL 2025	✓	Discrete Token	✗	—	Continuous
$\pi_{0.5}$ (Intelligence et al., 2025)	CoRL 2025	✓	Discrete Token	✗	—	Continuous
ThinkAct (Huang et al., 2025)	NeurIPS 2025	✓	Discrete Token	✗	—	Continuous
<i>VLAs with Visual CoT</i>						
CoT-VLA (Zhao et al., 2025)	CVPR 2025	✗	—	✓	Discrete Visual Tokens	Discrete
DreamVLA (Zhang et al., 2025b)	NeurIPS 2025	✗	—	✓	Discrete Visual Tokens	Continuous
UD-VLA (Chen et al., 2025b)	arXiv 2025	✗	—	✓	Discrete Visual Tokens	Discrete
VITA (Ma et al., 2025)	arXiv 2025	✗	—	✓	Discrete Visual Tokens	Discrete
<i>VLAs with Both Text and Visual CoT</i>						
UP-VLA (Zhang et al., 2025a)	ICML 2025	✓	Discrete Token	✓	Discrete Visual Tokens	Discrete
LaRA-VLA (Ours)	—	✓	Continuous Latent	✓	Continuous Latent	Continuous

130 2025; Lin et al., 2025b) since the introduction of RT-  
 131 2 (Zitkovich et al., 2023). Among these directions, CoT  
 132 reasoning at training time has proven particularly effective.  
 133 As summarized in Table 1, existing CoT-based VLA meth-  
 134 ods can be broadly categorized into textual CoT and visual  
 135 CoT. Textual CoT methods generate explicit textual reason-  
 136 ing traces to refine task instructions or extract motion- and  
 137 object-relevant information from visual observations (Za-  
 138 walski et al., 2025; Sun et al., 2025b; Huang et al., 2025). Vi-  
 139 sual CoT methods, in contrast, reconstruct or predict visual  
 140 observations conditioned on multimodal inputs, typically re-  
 141 lying on discrete visual tokens produced by VQ-VAE-based  
 142 representations (Zhao et al., 2025; Zhang et al., 2025b; Lv  
 143 et al., 2025). Despite their effectiveness, both paradigms  
 144 rely on discrete tokenized reasoning. Textual CoT incurs  
 145 long autoregressive generation chains with high inference  
 146 cost, while both textual and visual CoT exhibit a mismatch  
 147 with the inherently continuous perception and action spaces  
 148 in robotics. To address these limitations, we introduce la-  
 149 tent reasoning, in which textual CoT is replaced by com-  
 150 pact continuous latent variables learned via curriculum-style  
 151 training, and visual CoT is aligned with continuous visual  
 152 representations from the perception backbone.  
 153

## 2.2. Latent Reasoning in LLMs and VLMs

156 Explicit CoT can improve reasoning performance but often  
 157 incurs verbose intermediate outputs, high inference latency,  
 158 and a reliance on discrete tokens. These drawbacks have  
 159 motivated implicit or continuous reasoning in latent space,  
 160 which preserves multi-step computation without explicit  
 161 reasoning traces (Xu et al., 2025; Ruan et al., 2025). In  
 162 language models, latent reasoning is typically instantiated  
 163 through hidden states or continuous variables that function  
 164

as implicit thought. For instance, Coconut (Hao et al., 2024) shows that latent reasoning supports richer internal search while achieving a more favorable accuracy–efficiency trade-off than explicit CoT. SIM-CoT (Wei et al., 2025) identifies optimization instability when scaling implicit reasoning tokens and introduces supervised stabilization, while CoDi (Shen et al., 2025) distills explicit CoT into continuous latent representations via self-distillation. Latent reasoning has also been extended to vision–language models, where multimodal latents serve as internal reasoning states rather than explicit symbolic traces (Sun et al., 2025a; Pham & Ngo, 2025). To adapt latent reasoning to VLA, we adopt a staged training strategy that initializes structured reasoning with explicit CoT and progressively transfers it into latent space. Unlike prior methods that rely solely on answer supervision, we jointly ground latent reasoning in visual representations and action signals, enabling control-relevant reasoning that supports stable policy learning and efficient action generation without explicit CoT at inference time.

## 3. Method

In this section, we present the complete pipeline of our Latent Reasoning VLA (LaRA-VLA) framework. We first describe the construction of structured CoT datasets in Section 3.1. We then introduce the model architecture of LaRA-VLA and detail its training procedures in Sections 3.2 and 3.3, respectively.

### 3.1. Data Collection

Effective robotic manipulation requires the joint modeling of long-horizon subtask structure, spatial grounding of target objects, and motion-level reasoning for action execution.

165 However, existing data collection pipelines typically capture these components in isolation, leading to redundant or  
 166 incomplete supervision (e.g., exhaustive object-level bounding boxes in ECoT (Zawalski et al., 2025) or missing target  
 167 localization in Emma-x (Sun et al., 2025b)). To address  
 168 these limitations, we develop a fully automated annotation  
 169 pipeline following an anchor-first, generate-later paradigm  
 170 driven by semantic and temporal anchors. Semantic  
 171 anchors are extracted using Qwen3-VL (Bai et al., 2025a)  
 172 from the initial frame and task instruction to identify the  
 173 manipulated object, while temporal anchors are obtained  
 174 by segmenting robot trajectories into atomic manipulation  
 175 stages based on gripper state changes. Conditioned on these  
 176 anchors, Qwen3-VL generates concise subtask descriptions,  
 177 and open-vocabulary grounding with GroundingDINO (Liu  
 178 et al., 2024) and SAM3 (Carion et al., 2025) produces tempo-  
 179 rally consistent target object bounding boxes. Motion rea-  
 180 soning is derived from end-effector trajectories by comput-  
 181 ing goal-directed and local motions, which are discretized  
 182 into directional descriptors and incorporated into the CoT  
 183 annotations. Details are provided in Figure 9 and Appendix B.  
 184

185 Based on this pipeline, we construct two benchmark datasets  
 186 on simulated environments, LIBERO (Liu et al., 2023) and  
 187 SimplerEnv (Li et al., 2025), resulting in *LIBERO-LaRA*  
 188 and *Bridge-LaRA*, and further apply the same framework  
 189 to a long-horizon real-world robotic manipulation dataset  
 190 collected on physical hardware.

### 193 3.2. Model Architecture

195 We adopt Qwen3-VL (Bai et al., 2025a) as the backbone  
 196 VLM to leverage its strong built-in reasoning capability,  
 197 and directly inherit its image encoder to ensure consistent  
 198 visual representations throughout training. To predict visual  
 199 goal information, we introduce a dedicated <img\_next>  
 200 token to represent predicted visual latents, which enables  
 201 explicit supervision and alignment during early-stage latent  
 202 reasoning learning. For action prediction in Stages I and  
 203 II, we follow an autoregressive action token design simi-  
 204 lar to (Pertsch et al., 2025), allowing the model to learn  
 205 action generation jointly with latent reasoning in a stable  
 206 and efficient manner. In the final stage, we remove explicit  
 207 action token prediction and instead activate a dedicated ac-  
 208 tion expert, decoupling action generation from token-level  
 209 autoregressive decoding. Specifically, action generation is  
 210 performed by a 16-layer Diffusion Transformer composed of  
 211 alternating self-attention and cross-attention layers, which  
 212 conditions on the learned latent representations to produce  
 213 continuous action trajectories.

### 215 3.3. Training Procedures

217 **Stage I: Explicit CoT Fine-Tuning.** In the first stage,  
 218 we fine-tune the VLM on embodied datasets with explicit

219 CoT annotations constructed in Section 3.1. These annotations  
 220 provide structured reasoning supervision, including  
 221 task decomposition, movement reasoning, and target object  
 222 information, enabling the model to adapt to embodied ma-  
 223 nipulation tasks with explicit intermediate reasoning. Given  
 224 input images and a language instruction, the image encoder  
 225 first maps the visual observation to a sequence of visual  
 226 tokens, denoted as  $\mathbf{v}$ , while the instruction text is tokenized  
 227 into textual tokens, denoted as  $\mathbf{x}$ . The VLM jointly attends  
 228 to both visual and textual tokens and is trained to autoregres-  
 229 sively generate a sequence of CoT tokens. During this stage,  
 230 discrete ground-truth textual tokens are used to explicitly  
 231 supervise CoT generation via teacher forcing. The training  
 232 objective is defined as the negative log-likelihood of the  
 233 ground-truth CoT sequence:

$$\mathcal{L}_{\text{cot}} = - \sum_{t=1}^{T_{\text{cot}}} \log p_{\theta}(c_t | c_{<t}, \mathbf{v}, \mathbf{x}), \quad (1)$$

234 where  $\{c_t\}_{t=1}^{T_{\text{cot}}}$  is the ground-truth CoT token sequence,  $c_{<t}$   
 235 is the preceding CoT tokens, and  $p_{\theta}(\cdot)$  is the conditional  
 236 token distribution parameterized by the VLM.

237 In addition to textual supervision, we align visual reasoning  
 238 by predicting future visual latents. Let  $\mathbf{z}_{t+1}$  denote the vi-  
 239 sual latent representation of the next observation encoded by  
 240 the same visual encoder used for the input observations. The  
 241 VLM predicts this latent from the current context, yielding  
 242 the following alignment objective:

$$\mathcal{L}_{\text{vis}} = \|\hat{\mathbf{z}}_{t+1} - \mathbf{z}_{t+1}\|_1, \quad (2)$$

243 where  $\hat{\mathbf{z}}_{t+1}$  is the predicted visual latent and  $\mathbf{z}_{t+1}$  is obtained  
 244 by encoding the next observation. To stabilize visual latent  
 245 learning and prevent representation collapse, we follow prior  
 246 work (Chen et al., 2025a) and update the parameters used to  
 247 compute target visual latents using an exponential moving  
 248 average (EMA) of the online encoder parameters:

$$\bar{\theta}_v^t = \tau_v \bar{\theta}_v^{t-1} + (1 - \tau_v) \theta_v^t, \quad (3)$$

249 where  $\theta_v^t$  denotes the parameters of the online visual encoder  
 250 at iteration  $t$ ,  $\bar{\theta}_v^t$  denotes the corresponding EMA-averaged  
 251 parameters used to compute stable target visual latents, and  
 252  $\tau_v$  is the decay rate.

253 We further leverage the predicted future visual representa-  
 254 tions together with the preceding visual and textual context  
 255 to infer actions via an inverse dynamics model. Specifi-  
 256 cally, we employ an inverse dynamics function  $f(\mathbf{v}_t, \mathbf{v}_{t+1} |$   
 257  $\mathbf{x}, c) = \mathbf{a}_t$ , which estimates the action that induces the  
 258 transition between consecutive visual states conditioned  
 259 on the instruction and the intermediate reasoning step. To  
 260 efficiently propagate action information across latent reason-  
 261 ing steps, we adopt the fast recursive generation framework  
 262 of (Pertsch et al., 2025), which assigns coarse-grained action




Figure 2. Overview of LaRA-VLA. Training proceeds in three stages: (i) explicit CoT fine-tuning with aligned visual prediction latents and inverse-dynamics supervision for actions; (ii) a curriculum-based transition from explicit CoT to compact text latents, gradually reducing the number of text tokens while increasing reliance on latent reasoning, where the latent representations are also implicitly supervised by visual and action signals; and (iii) adaptation of latent-conditioned VLM features to an action expert for efficient action generation without explicit CoT at inference time.

semantics to all latent representations. Concretely, action tokens are trained using an autoregressive objective, similar to Equation 1, yielding the action-token loss  $\mathcal{L}_{\text{act-dis}}$ .

**Stage II: Curriculum-based Replacement of Discrete CoT Tokens.** In the second stage, we internalize explicit textual reasoning into the latent space through a curriculum-based training strategy. Embodied reasoning is modeled as a sequence of continuous latent states, and discrete CoT tokens are progressively replaced by latent representations during training. Formally, we adopt the same training objectives as in Stage I, including the textual CoT likelihood and the visual latent prediction loss. However, instead of supervising all reasoning steps with discrete ground-truth CoT tokens, we gradually mask out subsets of CoT tokens and replace them with learnable latent representations. The proportion of discrete CoT tokens decreases over training according to a predefined curriculum schedule, until the entire chain-of-thought is fully internalized in the latent space.

**Stage III: Action Generation via Flow Matching.** To adapt latent reasoning to continuous control, we train the action prediction module using a flow matching objective. Let  $\mathbf{a}_t$  denote the ground-truth action at time step  $t$ , and let  $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  be Gaussian noise. Following the flow matching formulation, we define a linear interpolation between noise and action as  $\mathbf{a}_\tau = (1 - \tau)\epsilon + \tau \mathbf{a}_t$ , where  $\tau \sim \mathcal{U}(0, 1)$ . The action expert predicts a velocity field  $v_{\theta}(\mathbf{a}_\tau, \tau | \mathbf{h}_t)$  conditioned on a multi-modal latent context  $\mathbf{h}_t$ . Specifically,  $\mathbf{h}_t$  aggregates the latent representation

of the current visual observation and language instruction, the intermediate text-based reasoning latent, and the predicted future visual latent produced by the VLM. Through the inverse dynamics supervision applied in earlier stages, this latent context already encodes coarse-grained action-relevant information. As a result, we do not introduce an additional action latent, and actions are generated directly from the shared multi-modal latent representation. The flow matching loss is defined as

$$\mathcal{L}_{\text{act-con}} = \mathbb{E}_{\mathbf{a}_t, \epsilon, \tau} \left[ \|v_{\theta_a}(\mathbf{a}_\tau, \tau | \mathbf{h}_t) - (\mathbf{a}_t - \epsilon)\|_2^2 \right], \quad (4)$$

where  $\mathbf{h}_t$  denotes the multi-modal latent context produced by the VLM at time step  $t$ .

**LaRA Attention Mechanism.** We introduce an attention mechanism tailored to our three-stage training paradigm, as illustrated in Figure 3. The model operates on multimodal tokens including text, current image, future image, and action tokens, with attention constraints explicitly regulating cross-token information flow. Here, text tokens serve as a unified abstraction that corresponds to language instructions and textual chain-of-thought in Stages I and II, and to text latents in Stages II and III. In Stages I and II, future image tokens attend causally to text and current image tokens, while interacting bidirectionally among themselves. Action tokens are generated autoregressively: each action token attends to all preceding text, current image, and future image tokens, as well as previously generated action tokens. In Stage III, action tokens are excluded from the attention

Table 2. Performance comparisons with state-of-the-art methods on LIBERO, grouped by different CoT paradigms.

CoT Type	Method	Spatial	Goal	Object	Long	Avg
No CoT	OpenVLA (Kim et al., 2025b)	84.7	88.4	79.2	53.7	76.5
	$\pi_0$ (Black et al., 2024)	96.8	98.8	95.8	85.2	94.2
	OpenVLA-OFT (Kim et al., 2025a)	97.6	98.4	97.9	94.5	97.1
Textual CoT	ThinkAct (Huang et al., 2025)	88.3	91.4	87.1	70.9	84.4
	MolmoAct (Lee et al., 2025)	87.0	95.4	87.6	77.2	86.6
	$\pi_{0.5}$ (Intelligence et al., 2025)	98.8	98.2	98.0	92.4	96.8
	DeepThinkVLA (Yin et al., 2025)	99.0	96.6	96.4	96.2	97.0
Visual CoT	CoT-VLA (Zhao et al., 2025)	87.5	91.6	87.6	69.0	81.1
	DreamVLA (Zhang et al., 2025b)	97.5	94.0	89.5	89.5	92.6
	F1 (Lv et al., 2025)	98.2	97.8	95.4	91.3	95.7
	UD-VLA (Chen et al., 2025b)	94.1	95.7	91.2	89.6	92.7
Latent CoT	Fast-ThinkAct (Huang et al., 2026)	92.0	97.2	90.2	79.4	89.7
	<b>LaRA-VLA (Ours)</b>	96.4	98.6	99.8	96.6	<b>97.9</b>




Figure 3. Attention mechanism used in LaRA-VLA.

computation, and the model is trained solely over text and vision tokens under the same attention constraints.

**Loss Function.** Our training objective is stage-dependent and follows a curriculum-style design. In Stage I, we jointly optimize the CoT supervision loss, the visual alignment loss, and the discrete action loss, i.e.,  $\mathcal{L}_{\text{cot}} + 0.1\mathcal{L}_{\text{vis}} + \mathcal{L}_{\text{act-dis}}$ , to initialize structured reasoning and action grounding. In Stage II, we progressively anneal the explicit CoT supervision loss  $\mathcal{L}_{\text{cot}}$  to zero. The model is lastly optimized using  $0.2\mathcal{L}_{\text{vis}} + \mathcal{L}_{\text{act-dis}}$ , which promotes latent-space reasoning while maintaining accurate action semantics. Finally, in Stage III, discrete action supervision is replaced by continuous action regression, and we optimize  $\mathcal{L}_{\text{act-con}}$  to enable efficient continuous action generation.

## 4. Experiments

We evaluate the effectiveness of LaRA-VLA and the overall system through a comprehensive set of experiments spanning both simulation benchmarks and real-world robotic manipulation tasks. Our experiments are designed to ad-

dress the following questions:

- How effective is LaRA-VLA compared to state-of-the-art methods in simulation benchmarks? (Section 4.1)
- How well does LaRA-VLA perform on **long-horizon** real-world manipulation tasks compared to state-of-the-art approaches? (Section 4.2)
- How effective are the latent reasoning components in LaRA-VLA, and what additional advantages does our approach offer? (Section 4.3)

### 4.1. Simulation Experiments

**Benchmarks and Datasets.** We conduct experiments on two widely used benchmarks, LIBERO (Liu et al., 2023) and SimplerEnv (Li et al., 2025). LIBERO consists of four task suites, Spatial, Goal, Object, and Long, each containing 10 single-arm manipulation tasks. We report success rates for each suite and the overall average over 50 rollouts per task. SimplerEnv evaluates real-to-sim generalization of robot manipulation policies trained on real-world data. We evaluate on WidowX robots across four tasks and report per-task success rates and the overall average over 24 rollouts per task. Based on these benchmarks, we construct two training datasets, LIBERO-LaRA and Bridge-LaRA, which are used to train LaRA-VLA.

**Baselines.** For LIBERO, we compare against a broad set of state-of-the-art VLA methods covering different CoT paradigms. No-CoT baselines include OpenVLA (Kim et al., 2025b),  $\pi_0$  (Black et al., 2024), and OpenVLA-OFT (Kim et al., 2025a). Textual CoT baselines include ThinkAct (Huang et al., 2025), MolmoAct (Lee et al., 2025),  $\pi_{0.5}$  (Intelligence et al., 2025), and DeepThinkVLA (Yin et al., 2025). Visual CoT methods include CoT-VLA (Zhao et al., 2025), DreamVLA (Zhang et al., 2025b), F1 (Lv et al., 2025), and UD-VLA (Chen et al., 2025b). We additionally compare with the latent CoT method Fast-ThinkAct (Huang

Table 3. Performance comparisons with state-of-the-art methods on SimplerEnv-WindowX, grouped by different CoT paradigms.

CoT Type	Method	Put Spoon	Put Carrot	Stack Block	Put Eggplant	Avg
No CoT	OpenVLA (Kim et al., 2025b)	0.0	0.0	0.0	4.1	1.0
	Octo (Ghosh et al., 2024)	47.2	9.7	4.2	56.9	29.5
	OpenVLA-OFT (Kim et al., 2025a)	12.5	4.2	8.3	37.5	39.6
	$\pi_0$ (Black et al., 2024)	29.1	0.0	16.7	62.5	40.1
	CogACT (Li et al., 2024)	71.7	50.8	15.0	67.5	51.3
Textual CoT	ThinkAct (Huang et al., 2025)	58.3	37.5	8.7	70.8	43.8
Visual CoT	F1 (Lv et al., 2025)	50.0	70.8	50.0	66.7	59.4
	UD-VLA (Chen et al., 2025b)	58.3	62.5	54.1	75.0	62.5
Latent CoT	<b>LaRA-VLA (Ours)</b>	95.8	62.5	25.0	91.7	<b>68.8</b>



Figure 4. Real-world setup of four long-horizon tasks.

et al., 2026). For SimplerEnv, we follow the standard evaluation protocol and compare against representative baselines. No-CoT methods include OpenVLA, Octo (Ghosh et al., 2024), OpenVLA-OFT,  $\pi_0$ , and CogACT (Li et al., 2024). We further include textual CoT (ThinkAct) and visual CoT (F1, UD-VLA) baselines.

**Results.** Tables 2 and 3 summarize quantitative results on the LIBERO and SimplerEnv-WidowX benchmarks. On LIBERO, LaRA-VLA achieves the best overall performance with an average success rate of 97.9%, including 99.8% on the Object suite and 96.6% on the Long suite, demonstrating strong object-centric reasoning and robustness in long-horizon manipulation. On SimplerEnv-WidowX, which evaluates real-to-sim generalization under diverse visual conditions, LaRA-VLA attains the highest average success rate of 68.8%, outperforming No-CoT, Textual CoT, and Visual CoT baselines. Across both benchmarks, LaRA-VLA consistently surpasses textual and visual CoT methods, indicating that latent reasoning provides more effective and stable guidance for action prediction and generalizes better than explicit CoT supervision.




Figure 5. Real-world results.

## 4.2. Real-world Experiments

**Real-World Setup.** As illustrated in Figure 4, our real-world setup uses an Agilex Cobot Magic wheeled platform equipped with three RGB-D cameras. We consider four categories of long-horizon manipulation tasks: placing all objects into a basket, sorting fruits into a basket, finding a block and placing it into a basket, and stacking two bowls. For data collection, we record 100 demonstration trajectories per task category at 30 Hz. During evaluation, each task is executed for 12 rollout trials. We compare our method against ACT (Zhao et al., 2023) and GR00T N1.5 (Bjorck et al., 2025) as baselines. Implementation details are provided in Appendix A.

**Results.** As shown in Figure 5, LaRA-VLA consistently outperforms ACT and GR00T N1.5 across all four long-horizon real-world manipulation tasks, achieving the highest average success rate. The improvements are especially pronounced on tasks requiring multi-stage reasoning and sustained temporal coordination, highlighting enhanced robustness to error accumulation over long horizons.

## 4.3. Analysis

**Ablation Study.** Table 4 shows that latent CoT supervision provides substantially larger gains than explicit textual CoT. While explicit textual CoT yields only marginal improvement over the no-CoT baseline, latent textual CoT leads

385  
386  
387  
388  
389  
390  
391  
392  
393  
394  
395  
396  
397  
398  
399  
400  
401  
402  
403  
404  
405  
406  
407  
408  
409  
410  
411  
412  
413  
414  
415  
416  
417  
418  
419  
420  
421  
422  
423  
424  
425  
426  
427  
428  
429  
430  
431  
432  
433  
434  
435  
436  
437  
438  
439

**Table 4.** Ablation study of different forms of CoT supervision on SimplerEnv. Text-CoT denotes explicit textual chain-of-thought, Latent Text-CoT denotes latent textual chain-of-thought, and Latent Vis-CoT denotes latent visual chain-of-thought.

Text-CoT	Latent Text-CoT	Latent Vis-CoT	SR (%)
✗	✗	✗	55.21
✓	✗	✗	58.33
✗	✓	✗	64.58
✗	✓	✓	<b>68.75</b>

to a significant increase in success rate, highlighting the effectiveness of internalized textual reasoning. Further incorporating latent visual CoT yields the best performance, as it both injects predictive visual information about future states and implicitly regularizes the latent textual CoT through multimodal alignment. Overall, these results indicate that structured latent CoT representations, especially when jointly learned across modalities, are more effective than explicit textual reasoning for embodied policy learning.

**Latent Collapse.** As shown in Figure 6, we observe no evidence of latent representation collapse. Latent tokens associated with different reasoning components form well-separated and semantically coherent clusters, demonstrating clear functional specialization rather than degeneration into uniform or uninformative representations. Moreover, latent representations of language instruction tokens (gray points) remain structured and occupy a distinct subspace from reasoning latents, indicating that latent CoT does not trivially reuse language embeddings. Together, these findings show that predictive supervision and action grounding provide sufficient inductive bias to maintain structured latent reasoning, even without explicit discrete chain-of-thought generation at inference time.

**Inference Efficiency.** As shown in Figure 7, LaRA-VLA significantly reduces inference latency, achieving 135 ms per rollout and outperforming all baselines by a large margin. Compared to explicit CoT methods, this yields up to a 90% reduction in inference time, demonstrating the efficiency benefits of latent reasoning without explicit CoT decoding.

## 5. Limitations

Although LaRA-VLA achieves fast inference and strong performance through latent chain-of-thought reasoning, several limitations remain and warrant further investigation. First, latent CoT representations may suffer from collapse in the absence of explicit supervision, where the semantics of latent tokens degenerate toward homogeneous representations, particularly as the number of latent tokens increases (Wei et al., 2025). To mitigate this risk, our current implementation restricts latent reasoning to a single token per step, which may limit expressiveness. Second, the training procedure is not maximally efficient. LaRA-VLA employs




Figure 6. Latent collapse analysis.




Figure 7. Inference time comparison on an NVIDIA A100 GPU.

a curriculum learning strategy that gradually replaces explicit CoT tokens with latent representations. As training progresses, the number of CoT-related tokens increases, resulting in higher training cost. Improving training efficiency while preserving stable latent reasoning remains an important direction for future work.

## 6. Conclusion

We presented LaRA-VLA, a latent reasoning framework for Vision–Language–Action models that internalizes chain-of-thought reasoning into continuous latent representations across both textual and visual modalities. Rather than generating long explicit CoT sequences at inference time, LaRA-VLA replaces them with compact textual CoT latents and employs a curriculum-based training strategy to progressively transfer explicit reasoning into latent space. Visual latents are aligned with continuous perceptual features encoded by a shared visual encoder and stabilized using an exponential moving average, providing implicit supervisory signals that guide the learning of textual CoT latents. Experiments on simulated benchmarks and long-horizon real-robot manipulation tasks demonstrate that LaRA-VLA achieves strong performance while significantly improving inference efficiency, supporting the view that structured reasoning for embodied agents can be effectively realized in latent space without explicit chain-of-thought generation.

440  
441 **Impact Statement**  
442

This paper aims to advance machine learning methods for robotic manipulation by improving the efficiency and scalability of reasoning in Vision–Language–Action models. The proposed approach introduces architectural and training innovations that enable effective reasoning without explicit chain-of-thought generation, potentially facilitating real-time robotic deployment. While such advances may benefit a wide range of robotic applications, including automation and assistive technologies, we do not identify any societal risks uniquely introduced by this work beyond those generally associated with robotic systems. As with related research, the broader impact will depend on how the technology is applied and governed in practice.

455  
456 **References**

Bai, S., Cai, Y., Chen, R., Chen, K., Chen, X., Cheng, Z., Deng, L., Ding, W., Gao, C., Ge, C., et al. Qwen3-vl technical report, 2025a. URL <https://arxiv.org/abs/2511.21631>, 2025a.

Bai, S., Zhou, W., Ding, P., Zhao, W., Wang, D., and Chen, B. Rethinking latent redundancy in behavior cloning: An information bottleneck approach for robot manipulation. In *Forty-second International Conference on Machine Learning*, 2025b.

Bjorck, J., Castañeda, F., Cherniadev, N., Da, X., Ding, R., Fan, L., Fang, Y., Fox, D., Hu, F., Huang, S., et al. Gr00t n1: An open foundation model for generalist humanoid robots. *arXiv preprint arXiv:2503.14734*, 2025.

Black, K., Brown, N., Driess, D., Esmail, A., Equi, M., Finn, C., Fusai, N., Groom, L., Hausman, K., Ichter, B., et al. pi\_0: A vision-language-action flow model for general robot control. *arXiv preprint arXiv:2410.24164*, 2024.

Carion, N., Gustafson, L., Hu, Y.-T., Debnath, S., Hu, R., Suris, D., Ryali, C., Alwala, K. V., Khedr, H., Huang, A., et al. Sam 3: Segment anything with concepts. *arXiv preprint arXiv:2511.16719*, 2025.

Chen, D., Shukor, M., Moutakanni, T., Chung, W., Yu, J., Kasarla, T., Bolourchi, A., LeCun, Y., and Fung, P. Vljepa: Joint embedding predictive architecture for vision-language. *arXiv preprint arXiv:2512.10942*, 2025a.

Chen, J., Song, W., Ding, P., Zhou, Z., Zhao, H., Tang, F., Wang, D., and Li, H. Unified diffusion vla: Vision-language-action model via joint discrete denoising diffusion process. *arXiv preprint arXiv:2511.01718*, 2025b.

Chen, W., Belkhale, S., Mirchandani, S., Mees, O., Driess, D., Pertsch, K., and Levine, S. Training strategies for efficient embodied reasoning. *arXiv preprint arXiv:2505.08243*, 2025c.

Deng, S., Yan, M., Wei, S., Ma, H., Yang, Y., Chen, J., Zhang, Z., Yang, T., Zhang, X., Zhang, W., et al. Graspvla: a grasping foundation model pre-trained on billion-scale synthetic action data. *Conference on Robot Learning*, 2025.

Ghosh, D., Walke, H. R., Pertsch, K., Black, K., Mees, O., Dasari, S., Hejna, J., Kreiman, T., Xu, C., Luo, J., et al. Octo: An open-source generalist robot policy. In *Robotics: Science and Systems*, 2024.

Hao, S., Sukhbaatar, S., Su, D., Li, X., Hu, Z., Weston, J., and Tian, Y. Training large language models to reason in a continuous latent space. *arXiv preprint arXiv:2412.06769*, 2024.

Huang, C.-P., Wu, Y.-H., Chen, M.-H., Wang, Y.-C. F., and Yang, F.-E. Thinkact: Vision-language-action reasoning via reinforced visual latent planning. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2025.

Huang, C.-P., Man, Y., Yu, Z., Chen, M.-H., Kautz, J., Wang, Y.-C. F., and Yang, F.-E. Fast-thinkact: Efficient vision-language-action reasoning via verbalizable latent planning. *arXiv preprint arXiv:2601.09708*, 2026.

Intelligence, P., Black, K., Brown, N., Darpinian, J., Dhabalnia, K., Driess, D., Esmail, A., Equi, M., Finn, C., Fusai, N., et al.  $\pi_{0.5}$ : A vision-language-action model with open-world generalization. In *Conference on Robot Learning*, 2025.

Kim, M. J., Finn, C., and Liang, P. Fine-tuning vision-language-action models: Optimizing speed and success. In *Robotics: Science and Systems*, 2025a.

Kim, M. J., Pertsch, K., Karamcheti, S., Xiao, T., Balakrishna, A., Nair, S., Rafailov, R., Foster, E. P., Sanketi, P. R., Vuong, Q., et al. Openvla: An open-source vision-language-action model. In *Conference on Robot Learning*, pp. 2679–2713. PMLR, 2025b.

Lee, J., Duan, J., Fang, H., Deng, Y., Li, B., Liu, S., Fang, B., Zhang, J., Wang, Y. R., Lee, S., et al. Molmoact: Action reasoning models that can reason in space. In *Workshop on Making Sense of Data in Robotics: Composition, Curation, and Interpretability at Scale at CoRL 2025*, 2025.

Li, Q., Liang, Y., Wang, Z., Luo, L., Chen, X., Liao, M., Wei, F., Deng, Y., Xu, S., Zhang, Y., et al. Cogact: A foundational vision-language-action model for synergizing cognition and action in robotic manipulation. *arXiv preprint arXiv:2411.19650*, 2024.

- 495 Li, X., Hsu, K., Gu, J., Mees, O., Pertsch, K., Walke, H. R.,  
 496 Fu, C., Lunawat, I., Sieh, I., Kirmani, S., et al. Evaluating  
 497 real-world robot manipulation policies in simulation. In  
 498 *Conference on Robot Learning*, pp. 3705–3728. PMLR,  
 499 2025.
- 500 Lin, F., Nai, R., Hu, Y., You, J., Zhao, J., and Gao, Y.  
 501 Onetwovla: A unified vision-language-action model with  
 502 adaptive reasoning. *arXiv preprint arXiv:2505.11917*,  
 503 2025a.
- 504 Lin, J., Taherin, A., Akbari, A., Akbari, A., Lu, L., Chen, G.,  
 505 Padir, T., Yang, X., Chen, W., Li, Y., et al. Vote: vision-  
 506 language-action optimization with trajectory ensemble  
 507 voting. *arXiv preprint arXiv:2507.05116*, 2025b.
- 508 Liu, B., Zhu, Y., Gao, C., Feng, Y., Liu, Q., Zhu, Y., and  
 509 Stone, P. Libero: Benchmarking knowledge transfer for  
 510 lifelong robot learning. *Advances in Neural Information  
 511 Processing Systems*, 36:44776–44791, 2023.
- 512 Liu, S., Zeng, Z., Ren, T., Li, F., Zhang, H., Yang, J., Jiang,  
 513 Q., Li, C., Yang, J., Su, H., et al. Grounding dino: Marry-  
 514 ing dino with grounded pre-training for open-set object  
 515 detection. In *European conference on computer vision*,  
 516 pp. 38–55. Springer, 2024.
- 517 Lv, Q., Kong, W., Li, H., Zeng, J., Qiu, Z., Qu, D., Song, H.,  
 518 Chen, Q., Deng, X., and Pang, J. F1: A vision-language-  
 519 action model bridging understanding and generation to  
 520 actions. *arXiv preprint arXiv:2509.06951*, 2025.
- 521 Ma, X., Xing, L., Zhang, H., Li, W., and Lu, S. Unify-  
 522 ing perception and action: A hybrid-modality pipeline  
 523 with implicit visual chain-of-thought for robotic action  
 524 generation. *arXiv preprint arXiv:2511.19859*, 2025.
- 525 Pertsch, K., Stachowicz, K., Ichter, B., Driess, D., Nair,  
 526 S., Vuong, Q., Mees, O., Finn, C., and Levine, S. Fast:  
 527 Efficient action tokenization for vision-language-action  
 528 models. *arXiv preprint arXiv:2501.09747*, 2025.
- 529 Pham, T.-H. and Ngo, C. Multimodal chain of continu-  
 530 ous thought for latent-space reasoning in vision-language  
 531 models. *arXiv preprint arXiv:2508.12587*, 2025.
- 532 Ruan, Y., Band, N., Maddison, C. J., and Hashimoto, T.  
 533 Reasoning to learn from latent thoughts. *arXiv preprint  
 534 arXiv:2503.18866*, 2025.
- 535 Shen, Z., Yan, H., Zhang, L., Hu, Z., Du, Y., and He, Y. Codi:  
 536 Compressing chain-of-thought into continuous space via  
 537 self-distillation. *arXiv preprint arXiv:2502.21074*, 2025.
- 538 Sridhar, K., Dutta, S., Jayaraman, D., and Lee, I. Ricl:  
 539 Adding in-context adaptability to pre-trained vision-  
 540 language-action models. In *Conference on Robot Learn-  
 541 ing*, pp. 5022–5038. PMLR, 2025.
- 542 Sun, G., Hua, H., Wang, J., Luo, J., Dianat, S., Rabbani, M.,  
 543 Rao, R., and Tao, Z. Latent chain-of-thought for visual  
 544 reasoning. *arXiv preprint arXiv:2510.23925*, 2025a.
- 545 Sun, Q., Hong, P., Pala, T. D., Toh, V., Tan, U.-X., Ghosal,  
 546 D., and Poria, S. Emma-x: An embodied multimodal  
 547 action model with grounded chain of thought and look-  
 548 ahead spatial reasoning. In *Proceedings of the 63rd  
 549 Annual Meeting of the Association for Computational  
 550 Linguistics (Volume 1: Long Papers)*, pp. 14199–14214,  
 551 2025b.
- 552 Wei, X., Liu, X., Zang, Y., Dong, X., Cao, Y., Wang, J., Qiu,  
 553 X., and Lin, D. Sim-cot: Supervised implicit chain-of-  
 554 thought. *arXiv preprint arXiv:2509.20317*, 2025.
- 555 Xu, Y., Guo, X., Zeng, Z., and Miao, C. Softcot: Soft  
 556 chain-of-thought for efficient reasoning with llms. *arXiv  
 557 preprint arXiv:2502.12134*, 2025.
- 558 Yin, C., Lin, Y., Xu, W., Tam, S., Zeng, X., Liu, Z., and  
 559 Yin, Z. Deepthinkvla: Enhancing reasoning capabili-  
 560 ty of vision-language-action models. *arXiv preprint  
 561 arXiv:2511.15669*, 2025.
- 562 Zawalski, M., Chen, W., Pertsch, K., Mees, O., Finn, C.,  
 563 and Levine, S. Robotic control via embodied chain-of-  
 564 thought reasoning. In *Conference on Robot Learning*, pp.  
 565 3157–3181. PMLR, 2025.
- 566 Zhang, J., Guo, Y., Hu, Y., Chen, X., Zhu, X., and Chen, J.  
 567 Up-vla: A unified understanding and prediction model  
 568 for embodied agent. In *Forty-second International Con-  
 569 ference on Machine Learning*, 2025a.
- 570 Zhang, W., Liu, H., Qi, Z., Wang, Y., Yu, X., Zhang, J.,  
 571 Dong, R., He, J., Wang, H., Zhang, Z., et al. Dreamvla:  
 572 A vision-language-action model dreamed with compre-  
 573 hensive world knowledge. In *The Thirty-ninth Annual  
 574 Conference on Neural Information Processing Systems*,  
 575 2025b.
- 576 Zhao, Q., Lu, Y., Kim, M. J., Fu, Z., Zhang, Z., Wu, Y.,  
 577 Li, Z., Ma, Q., Han, S., Finn, C., et al. Cot-vla: Visual  
 578 chain-of-thought reasoning for vision-language-action  
 579 models. In *Proceedings of the Computer Vision and  
 580 Pattern Recognition Conference*, pp. 1702–1713, 2025.
- 581 Zhao, T., Kumar, V., Levine, S., and Finn, C. Learning fine-  
 582 grained bimanual manipulation with low-cost hardware.  
 583 *Robotics: Science and Systems XIX*, 2023.
- 584 Zitkovich, B., Yu, T., Xu, S., Xu, P., Xiao, T., Xia, F.,  
 585 Wu, J., Wohlhart, P., Welker, S., Wahid, A., et al. Rt-2:  
 586 Vision-language-action models transfer web knowledge  
 587 to robotic control. In *Conference on Robot Learning*, pp.  
 588 2165–2183. PMLR, 2023.

## 550 A. Implementation Details

### 551 A.1. Implementation Details of LaRA-VLA

553 **Training Paradigm.** We adopt the progressive three-stage training paradigm across all experiments. In **Stage I**, the  
 554 input sequence encompasses explicit Chain-of-Thought (CoT) annotations, `<img_next>` tokens for next-frame feature  
 555 prediction, and action tokens tokenized via fast (Pertsch et al., 2025) tokenizer. In **Stage II**, we facilitate the learning of  
 556 implicit reasoning by substituting concrete CoT steps with `<thinking>` tokens. Finally, in **Stage III**, we discard explicit  
 557 CoT supervision. The specific input formats for each stage are illustrated in Figure 8.  
 558

```

559 Stage I: Explicit Supervision
560 'Place the potato inside the bowl. @ Subtask: carry the potato toward the bowl.
561 BBox: [0.5664 0.5898 0.6953 0.6641]. Reasoning: the robot is closing the gripper.
562 <img_next> <img_next> <img_next> <img_next> <img_next> <img_next>
563 <img_next> <img_next> <img_next> <img_next> <img_next> <img_next>
564 <img_next> <img_next> <img_next> <img_next> <img_next> <img_next>
565 <img_next> <img_next> <img_next> <img_next> <img_next> <img_next>
566 Stage II: Mixed Strategy (Transition)
567 One thinking token:
568 'Place the potato inside the bowl. @ <|start_of_thinking|> <|thinking|>
569 <|end_of_thinking|> BBox: [0.5664 0.5898 0.6953 0.6641]. Reasoning: the robot is
570 closing the gripper. <img_next> <img_next> <img_next> <img_next> <img_next>
571 <img_next> <img_next> <img_next> <img_next> <img_next> <img_next>
572 <img_next> <img_next> <img_next> <img_next> <img_next> <img_next>
573 Two thinking tokens:
574 'Place the potato inside the bowl. @ <|start_of_thinking|> <|thinking|> <|thinking|>
575 <|end_of_thinking|> Reasoning: the robot is closing the gripper. <img_next> <img_next>
576 <img_next> <img_next> <img_next> <img_next> <img_next> <img_next>
577 <img_next> <img_next> <img_next> <img_next> <img_next> <img_next>
578 <img_next> <img_next> <img_next> <img_next> <img_next> <img_next>
579 Three thinking tokens:
580 'Place the potato inside the bowl. @ <|start_of_thinking|> <|thinking|> <|thinking|>
581 <|thinking|> <|end_of_thinking|> <img_next> <img_next> <img_next> <img_next>
582 <img_next> <img_next> <img_next> <img_next> <img_next> <img_next>
583 <img_next> <img_next> <img_next> <img_next> <img_next> <img_next>
584 Stage III: Fully Latent Reasoning
585 'Place the potato inside the bowl. @ <|start_of_thinking|> <|thinking|> <|thinking|>
586 <|thinking|> <|end_of_thinking|> <img_next> <img_next> <img_next> <img_next>
587 <img_next> <img_next> <img_next> <img_next> <img_next> <img_next>
588 <img_next> <img_next> <img_next> <img_next> <img_next> <img_next>
589 <img_next> <img_next> <img_next> <img_next> <img_next> <img_next>
590 <img_next> <img_next> <img_next> <img_next> <img_next> <img_next>
591
```

585 586 *Figure 8. Example of Training Data Formats.*

587 **Training Config.** Most hyperparameters are shared across evaluation settings, with only a small subset adapted to each  
 588 scenario. In particular, the action horizon is set to 16 for the Bridge dataset, 8 for the LIBERO benchmark, and 25 for  
 589 real-world experiments. Detailed hyperparameter configurations for the LIBERO, Bridge, and real-world evaluations are  
 590 summarized in Table 5. All models are trained using 8 NVIDIA H100 GPUs.  
 591

### 592 A.2. Implementation Details Baselines in Real-world Experiments

593 **ACT (Zhao et al., 2023).** We instantiate ACT using the LeRobot implementation, configured with a chunk size of  $K = 50$   
 594 and  $n_{action} = 50$  action steps per chunk. The perception stack follows the original design, using a ResNet-18 vision backbone  
 595 pretrained on ImageNet and mean-std normalization for visual, state, and action inputs. The transformer backbone employs  
 596 a 4-layer encoder and a 1-layer decoder with model dimension 512, 8 attention heads, and a 3200-dimensional feedforward  
 597 network with ReLU activation and dropout 0.1. Following prior work, we enable the VAE module with a 32-dimensional  
 598 latent space, 4 encoder layers, and a KL weight of 10.0. Training is performed for 40k gradient steps with batch size 100 on  
 599 a single H100 GPU, using Adam with learning rate  $1 \times 10^{-5}$ , weight decay  $1 \times 10^{-4}$ , and the same learning rate for the  
 600 backbone. Note that we adopt the LeRobot fix for the ACT decoder, which uses a single effective decoder layer consistent  
 601 with the original implementation.  
 602

603 **GR00T N1.5 (Bjorck et al., 2025).** We use the original GR00T N1.5 implementation with its default architecture and a  
 604

Table 5. Hyperparameter settings across LIBERO, SimplerEnv, and Real Robot experiments.

Hyperparameters	Libero			SimplerEnv			Real Robot		
	Stage I	Stage II	Stage III	Stage I	Stage II	Stage III	Stage I	Stage II	Stage III
<i>Learning Rates</i>									
VLM LR	$1 \times 10^{-5}$	$1.3 \times 10^{-5}$	$1 \times 10^{-5}$	$1 \times 10^{-5}$	$1 \times 10^{-5}$				
DiT LR	$1 \times 10^{-4}$	$1.3 \times 10^{-4}$	$1 \times 10^{-4}$	$1 \times 10^{-4}$	$1 \times 10^{-4}$				
<i>Optimization Config</i>									
Action Horizon	8	8	8	16	16	16	25	25	25
Training Steps	5k	2k/2k/2k	40k	10k	5k/5k/10k	60k	5k	2k/2k/2k	10k
Batch Size	12	16	16	12	16	16	12	16	16
Optimizer	AdamW	AdamW	AdamW	AdamW	AdamW	AdamW	AdamW	AdamW	AdamW
LR Scheduler	Cosine	Cosine	Cosine	Cosine	Cosine	Cosine	Cosine	Cosine	Cosine
Warm-up Ratio	0.1	0.1	0.1	0.1	0.1	0.1	0.1	0.1	0.1
<i>Loss Weights</i>									
Action Token Loss	1.0	1.0	–	1.0	1.0	–	1.0	1.0	–
Image Next Loss	0.1	0.2	–	0.1	0.2	–	0.1	0.2	–
CoT Loss	1.0	1.0	–	1.0	1.0	–	1.0	1.0	–
DiT Loss	–	–	1.0	–	–	1.0	–	–	1.0

continuous-action flow matching head. For real-robot experiments, we train GR00T with an action chunk size of  $K = 25$  and batch size 128 on a single H100 GPU. We follow the recommended optimization hyperparameters, using AdamW with learning rate  $1 \times 10^{-4}$ , weight decay  $1 \times 10^{-5}$ , and a warmup ratio of 0.05 of the total training steps. Consistent with the original setup, we freeze both the language model backbone and the vision tower, and fine-tune only the projector and diffusion policy head.

## B. Details of Data Pipeline

We argue that effective robotic manipulation relies on the organic integration of three tightly coupled components: subtask analysis, spatial grounding of target objects, and directional motion reasoning for the manipulator. Subtask decomposition enables long-horizon reasoning, spatial information localizes task-relevant objects, and motion reasoning translates high-level intent into executable control signals. However, existing data collection pipelines fail to jointly capture these components in a compact and task-centric manner. For example, ECoT (Zawalski et al., 2025) annotates bounding boxes for all scene objects, leading to highly redundant supervision, while Emma-x (Sun et al., 2025b) lacks explicit target localization. To address this gap, we construct a fully automated annotation pipeline without human intervention, following an anchor-first, generate-later paradigm driven by semantic and temporal anchors, as shown in Figure 9.

**Subtask Annotation.** We extract semantic anchors using Qwen3-VL (Bai et al., 2025a) to identify the manipulated object from the first-frame image and the task instruction, which provides a textual reference for subsequent visual grounding. The prompt used for object identification is illustrated in Figure 10. Temporal anchors are obtained by segmenting robot trajectories into atomic manipulation stages (e.g., pre-grasp, grasp, move, and release) based on changes in the gripper state, with segment boundaries treated as keyframes. Conditioned on the instruction and the corresponding keyframes, Qwen3-VL generates high-level subtask descriptions for each segment. The prompt used for Subtask Annotation is illustrated in Figure 11.

**Target Object Bounding Boxes.** Guided by the semantic anchors, we perform open-vocabulary spatial grounding using GroundingDINO (Liu et al., 2024) and SAM3 (Carion et al., 2025) to obtain temporally consistent 2D bounding box trajectories of the target object. To be noted, for the Bridge dataset, we improve the robustness of our bounding box annotations by employing a multi-frame ensemble approach. We prompt SAM using GroundingDINO detections from 5 uniformly sampled anchor frames. The final trajectory is obtained by filtering spatial outliers and selecting the highest-confidence sequence. Furthermore, linear interpolation is applied to address tracking discontinuities, ensuring dense frame-wise supervision.

**Motion Reasoning.** Motion reasoning is derived from end-effector state trajectories by computing both global motion toward the segment goal and local instantaneous motion. These motion vectors are further mapped to directional descriptors




Figure 9. Data collection pipeline.

and incorporated into the CoT annotations.

## C. Analysis of Real-world Experiments

Table 6 provides a subtask-level breakdown of real-world manipulation performance, revealing clear differences in temporal dependency across tasks. For tasks such as Put All Objects into the Basket, Put All Fruits into the Basket, and Stack Two Bowls, the two subtasks are largely decoupled and can be completed independently. This is reflected by the fact that failures in one subtask do not necessarily propagate to the other, and improvements in overall success rate are primarily driven by independent gains in individual subtasks. In contrast, Find the Block and Place It in the Basket exhibits strong sequential dependency between subtasks. Successful completion of the second subtask (placing the block) is contingent on correctly executing the first subtask (finding the block). As a result, errors in the initial subtask directly limit the achievable success rate of the subsequent subtask, leading to a tightly coupled failure mode. This structural dependency explains why improvements on this task require coherent reasoning and action execution across subtasks, rather than isolated subtask-level optimization. Notably, LaRA-VLA demonstrates consistent gains over GR00T N1.5 on this tightly coupled task, suggesting that its latent reasoning mechanism is particularly effective in maintaining cross-subtask coherence. By contrast, for tasks with weak subtask coupling, performance improvements can be largely attributed to better local action execution, and the benefits of reasoning primarily manifest as incremental gains.

Table 6. Subtask-level and overall success rates (%) on real-world robot tasks.

Method	Put All Objects into the Basket			Put All Fruits into the Basket			Find the Block and Place It in the Basket			Stack Two Bowls		
	Subtask 1	Subtask 2	Overall	Subtask 1	Subtask 2	Overall	Subtask 1	Subtask 2	Overall	Subtask 1	Subtask 2	Overall
ACT	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
GR00T N1.5	50.0	50.0	<b>50.0</b>	50.0	83.3	33.3	100.0	16.7	16.7	91.7	100.0	91.7
LaRA-VLA	50.0	41.7	41.6	66.7	75.0	<b>50.0</b>	100.0	<b>33.3</b>	<b>33.3</b>	100.0	100.0	<b>100.0</b>

```

715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
    ...
You are helping a robot understand which object to manipulate.
You will receive:
1) A pre-grasp image.
2) The task instruction the robot must follow.
Instruction: ``{instruction}``
Your goal is to find the SINGLE primary object that the robot must interact with.
- Read the instruction carefully to identify the nouns that describe the object(s) to be manipulated.
- Prioritize what the instruction explicitly requests; ignore other objects unless they are essential context.
- Ignore objects that the robot does not need to touch.
- Describe the object with a short noun phrase (preferably including color if visible) but never mention any location context (e.g., no ``on stove'', ``in pot'', ``on left'').
- The description must be  $\leq 3$  words, all lowercase, and purely object attributes.
- List any other important objects mentioned in the instruction under secondary_objects (also  $\leq 3$  words, lowercase, no location mentions).
Output STRICT JSON with the following fields (no comments, no extra text):
``manipulated_object``: string describing the main object ( $\leq 3$  words, lowercase, no location).
``secondary_objects``: array of strings for other relevant instruction objects (each  $\leq 3$  words, lowercase, no location).
If you are uncertain or there is no target object, set ``manipulated_object`` to the noun mentioned in the instruction and secondary_objects to an empty array.
JSON:
...

```

*Figure 10.* Prompt for Object Identification.

```

770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785 You are a robot manipulation expert.
786 Your goal is to describe what the robot is doing in this phase of a pick-and-place
787 task.

788 Global instruction: ``{instruction}``
789 Internal segment label (for your reference only): ``{segment_type}```

790 You are given two keyframes from this phase: one near the beginning and one near the
791 end.
792 From the instruction and the images, you must:
793 1. Identify the single main object the robot is manipulating. If the object cannot be
    identified, return ``unknown``.
794 2. Describe where this object is located or how it is situated in the scene
    (scene_context), e.g., ``on the table'', ``in the basket'', or ``in the box''.
795 3. Write a short natural-language subtask description for the current action phase,
    following the rules below:
796     - The subtask MUST be a concrete action phrase (e.g., ``reach toward the blue cup'',
        ``carry the spoon toward the bowl'').
797     - DO NOT simply repeat generic labels such as ``move_to_object'', ``move_to_goal'',
        ``grasp_object'', or ``place_object''.
798     - The subtask should explicitly include the object name. If the object cannot be
        identified, return the subtask as ``manipulate the object''.

800 Return STRICT JSON only (no extra text) with fields:
801 {
802     ``object_name'': ``string (e.g., 'blue cup')'',
803     ``scene_context'': ``string (e.g., 'inside the metal basket')'',
804     ``subtask'': ``string (natural language, must include the object)``
805 }
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824

```

*Figure 11.* Prompt for Subtask Description Generation.